



# Nuvolaris Trainings

## Nuvolaris Operator

Part 1: Introducing Kubernetes  
Operators

<https://www.nuvolaris.io>

# Agenda (Part 1)


- The Nuvolaris Project
- Development Environment
- Kubernetes 101
- Custom Resource Definitions
- Kustomize

# What is Kubernetes ?

- In theory, an **orchestrator**
  - also Windows, originally, was just a **GUI** on top of DOS
- In practice, an **Operating System** for the cloud

## What is Nuvolaris?

- a Serverless **distribution** for Kubernetes
- *Linux* : **RedHat** = *Kubernetes* : **Nuvolaris**



# Nuvolaris Architecture

# Kubernetes Operators

- It is a **pattern** that is becoming commonplace
  - There is *NOT* a specific API that you implement
  - You have to use the *Kubernetes API* anyway
- You define your own Resource
  - Defining new resources as **CRD** Custom Resource Definitions
  - Creating instances conforming to the CRD
    - that describes the *desired state*
  - **Writing code that brings the system to this state**

# Operator Frameworks

- **Operator Framework:** ansible/helm/go
- **Kudo:** a declarative, yaml based framework
- **Metacontroller:** generic, with hooks in any languages
- **Shell-operator:** write operators in bash
- **Kubebuilder:** Go based operator
- **Kopf:** Python based Operator

also exists Java, Rust, Elixir, Javascript based operator frameworks

# Dev Environment

# VSCode-based Development Environment

- Clone the repositories (multiple and linked)

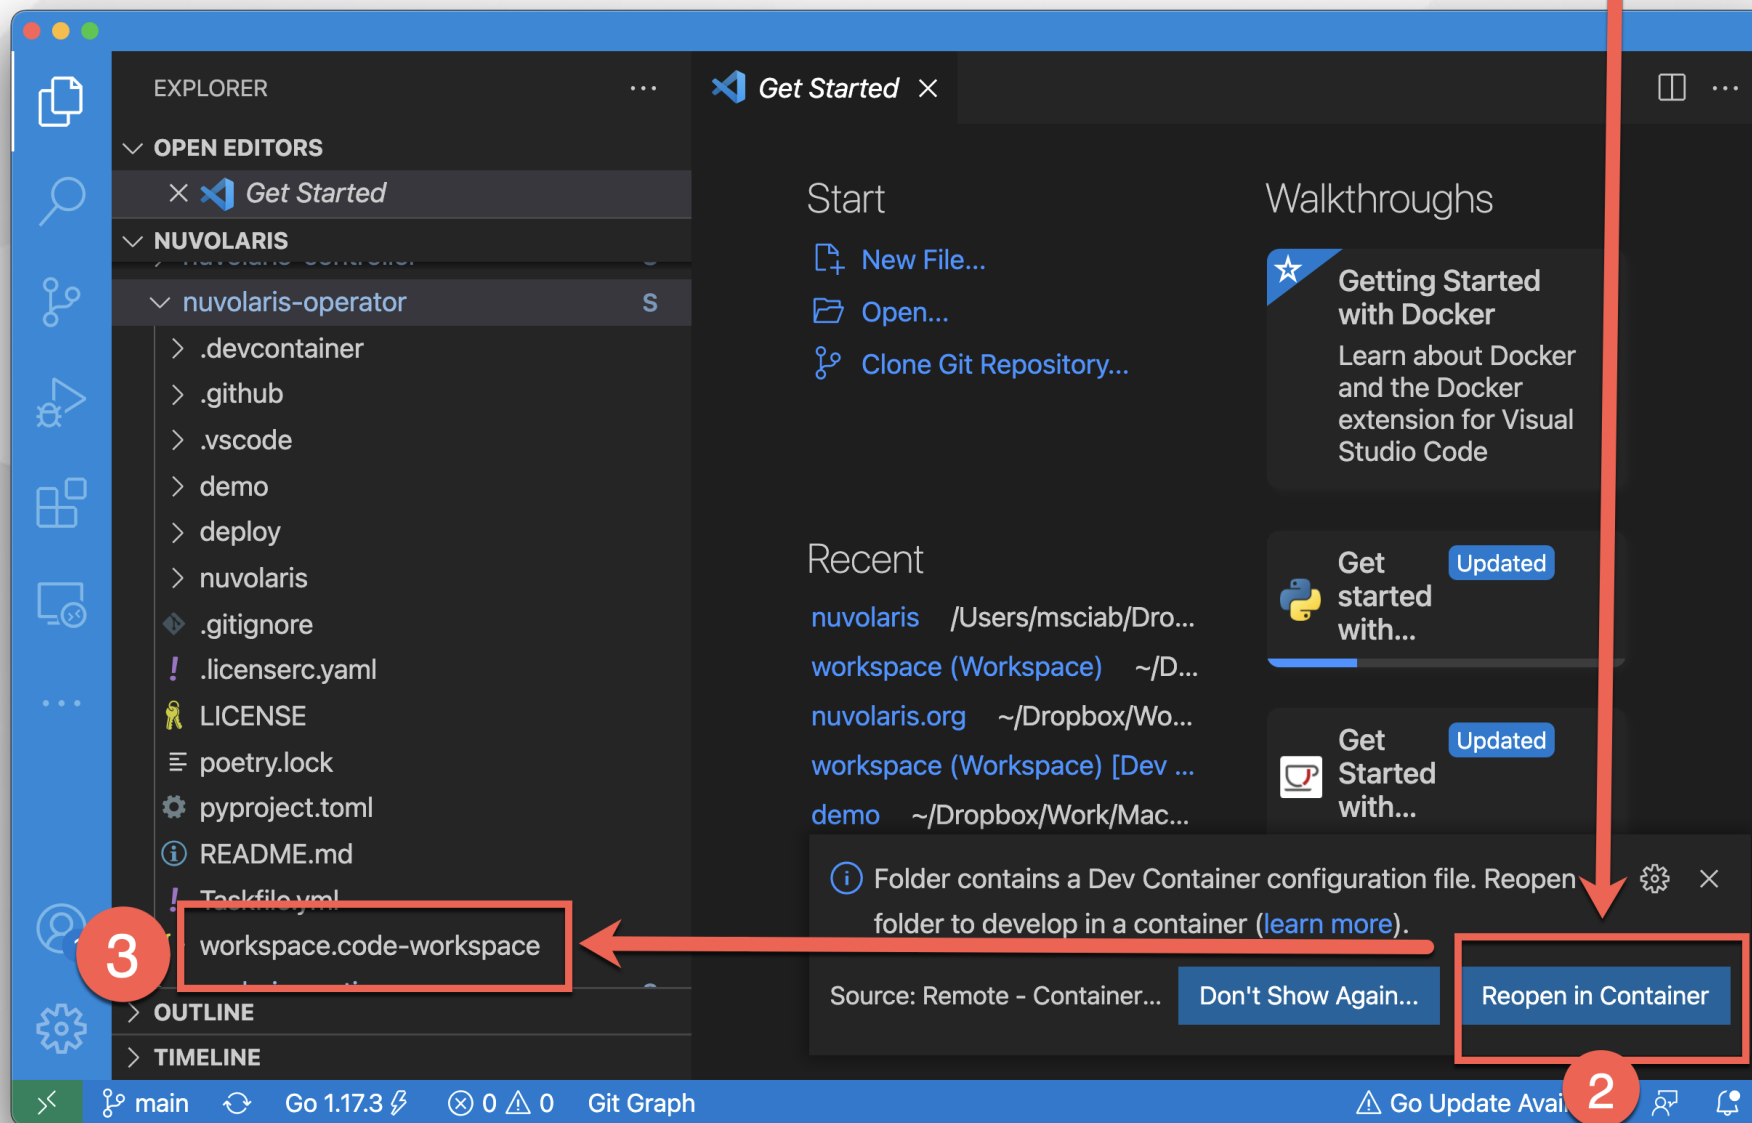
```
git clone https://github.com/nuvolaris/nuvolaris  
--recurse-submodules
```

- do not forget `--recurse-submodules`
- Open the folder `nuvolaris` with VSCode:
  - Command Line: `code nuvolaris`
- Open the workspaces in subfolders: `workspace.code-workspace`



1

```
$ git clone https://github.com/nuvolaris/nuvolaris --recurse-submodules  
$ code nuvolaris
```



# Test Nuvolaris Operator

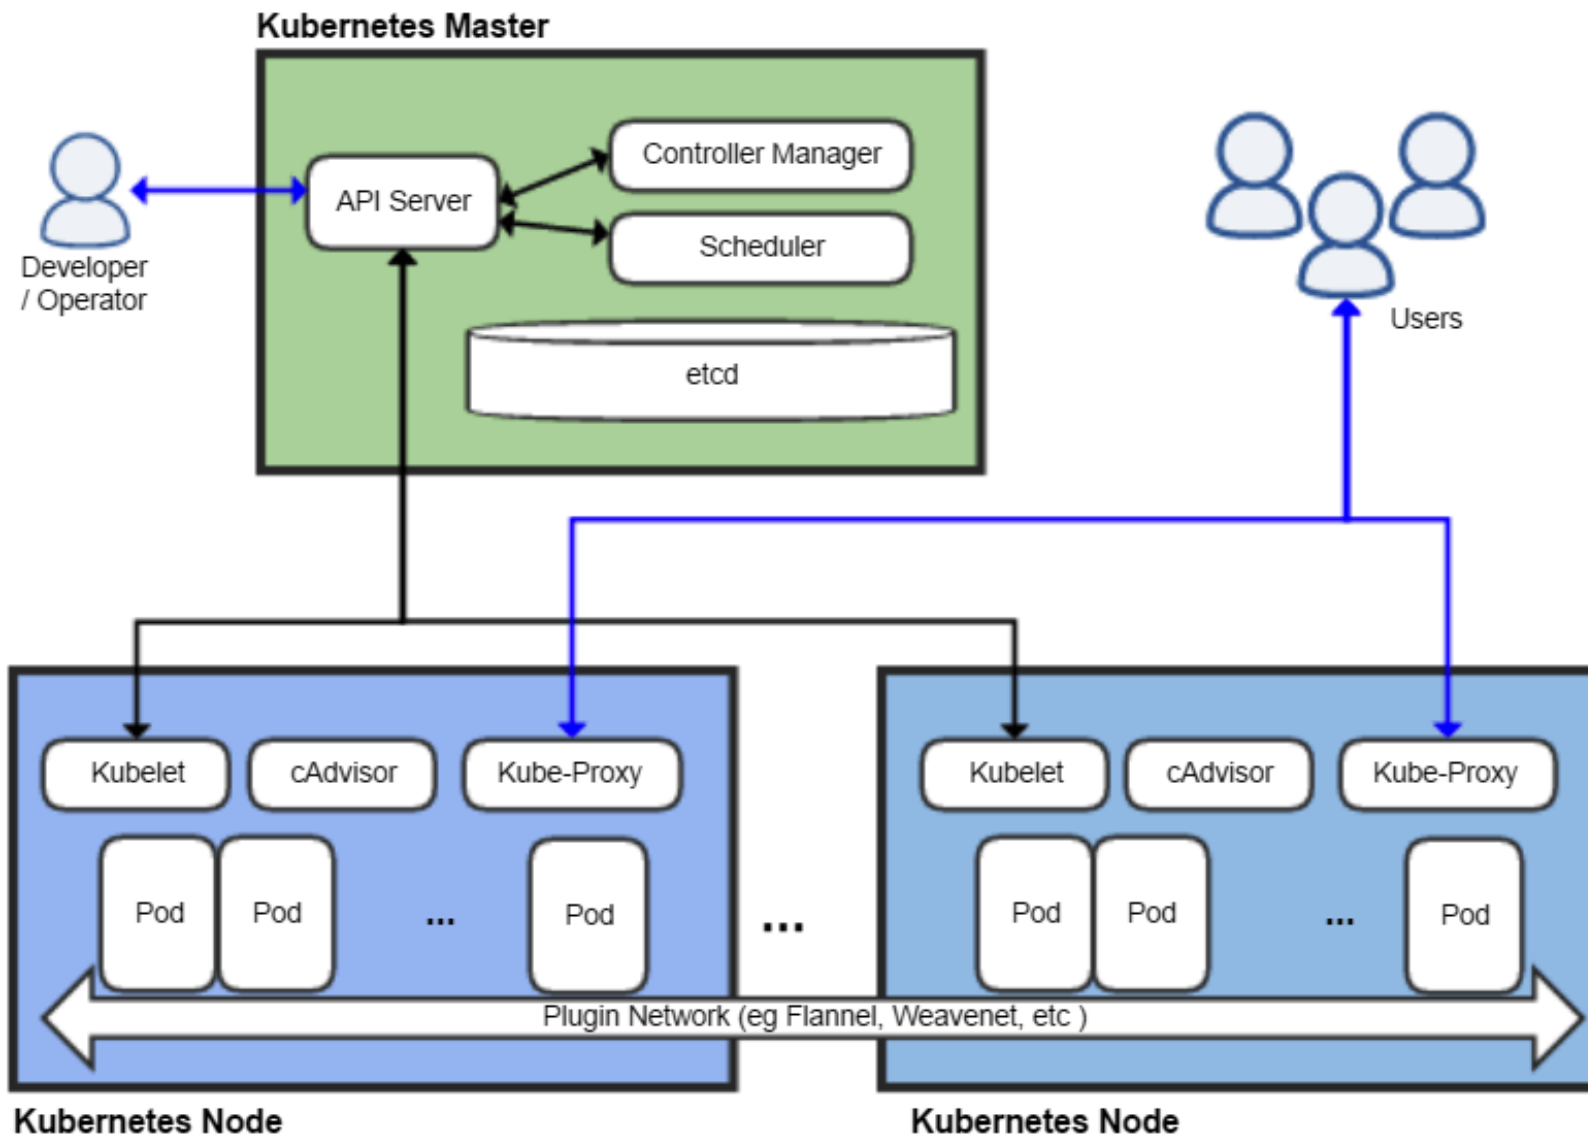
- Open `nuvolaris-operator/workspace.code-workspace`
- `kubectl get nodes`

NAME	STATUS	ROLES	AGE	VERSION
nuvolaris-control-plane	Ready	control-plane,master	41m	v1.21.1
nuvolaris-worker	Ready	<none>	41m	v1.21.1

- Test

```
cd tests
task deploy
task wsk
```

# Kubernetes 101



# Kubernetes Architecture

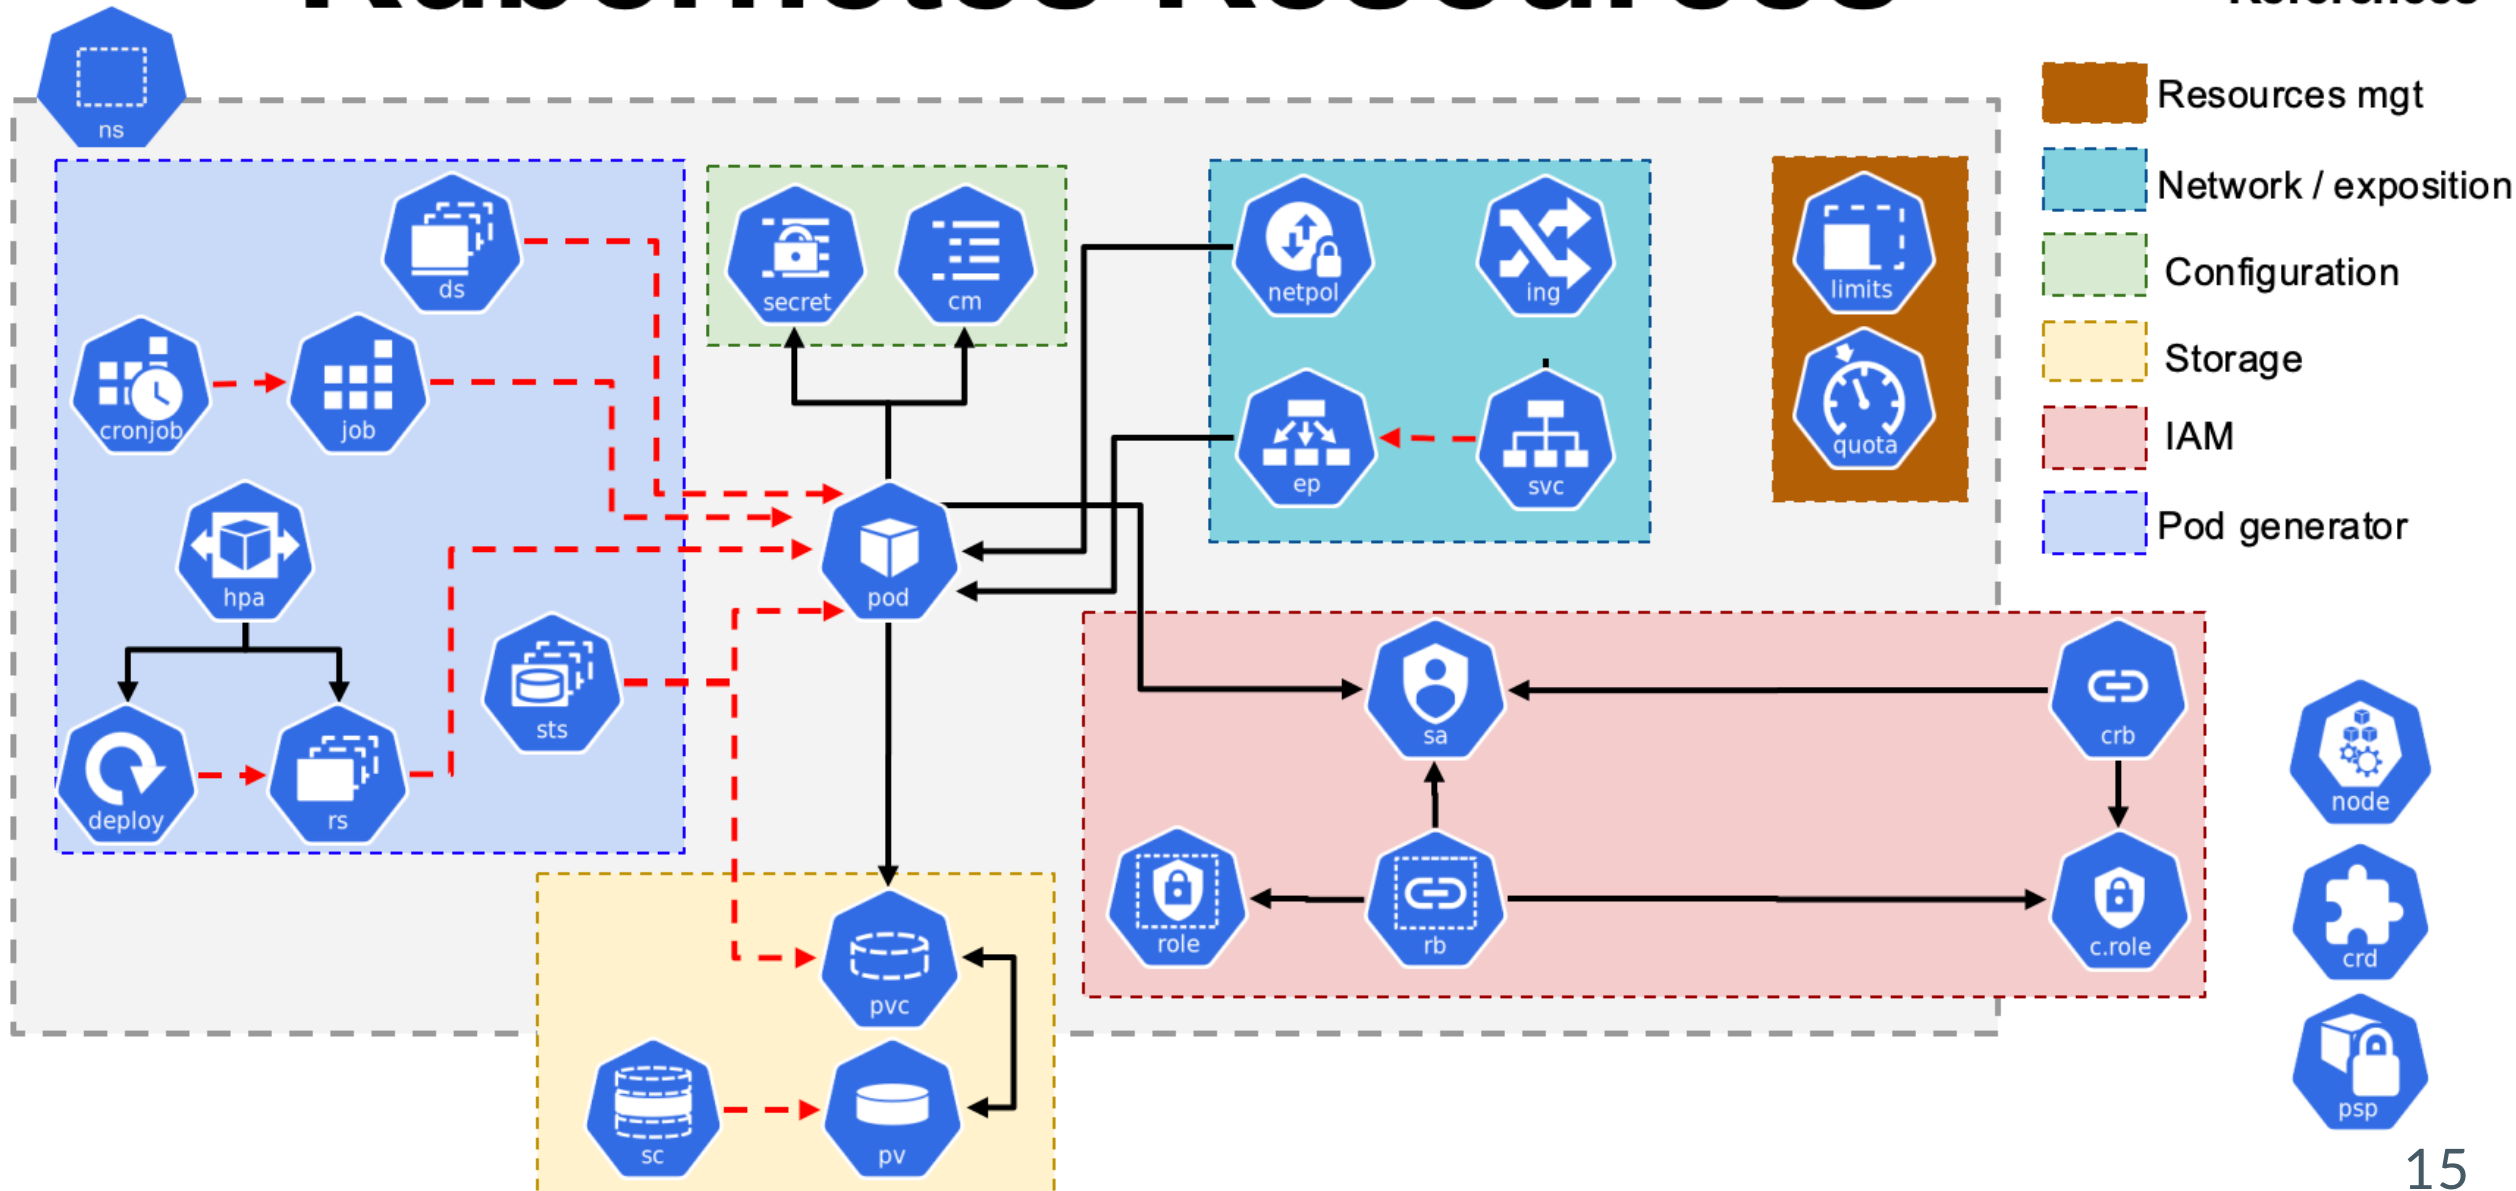
# Kubernetes **kubectl** Commands

```
# checking nodes and namespaces
kubectl get nodes
kubectl get ns
kubectl create ns demo
kubectl get ns
# default namespace
kubectl config set-context --current --namespace demo
# various resources
kubectl get pod
kubectl get deploy
kubectl get svc
```

# Kubernetes Descriptors Concepts

- Kubernetes is declarative:
  - You **describe** what you want to get by the system
  - Kubernetes will bring the system to the desired state
- You declare what you want with *descriptors* in YAML
  - those descriptors are in YAML format
    - actually, they are internally JSON files
    - YAML is really syntax sugar for JSON
- Kubernetes brings the system to what you asked
  - ... **if it is possible** ...

# Kubernetes Resources



# Structure of a Kubernetes resource

- Common: Header and Metadata

```
apiVersion: v1           # resources are versioned AND grouped
kind: Pod                # each resource has a `kind`
metadata:
  name: demo-pod         # name of the resources
  namespace: demo        # grouped in a namespace
  labels:                # used to locate resources
    app: demo            # formact key=value
```

- **spec**: changes according to the kind
- **status**: maintained by the system



# Simple Descriptor: a Pod

- A pod is a *set* of containers
  - the closest thing to `docker run`

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-pod
  namespace: demo
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```

# Deploy Pod

```
cat demo-pod.yaml  
kubectl apply -f demo-pod.yaml  
kubectl get pod  
kubectl delete pod demo-pod
```

# Nested Descriptor: a Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: demo-deploy
```

Templatized, repeat the template using labels

```
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
```

# Deployment template

It creates `replica` times the pods specified in the template

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx
        ports:
          - containerPort: 80
```

# Deploy a Deployment

```
cat demo-deployment.yaml
kubectl apply -f demo-deployment.yaml
kubectl get deploy
kubectl get pod
kubectl delete -f demo-deployment.yaml
kubectl get deploy
kubectl get pod
```

# Kubernetes CRD

# Kubernetes Controllers

- Deployment, DaemonSet, StatefulSet

```
[~]$ kubectl get deploy
NAME                READY    UP-TO-DATE    AVAILABLE    AGE
nginx-deployment    3/3      3              3             21m
[~]$ kubectl get po
NAME                                READY    STATUS    RESTARTS    AGE
nginx-deployment-66b6c48dd5-4dpl2    1/1      Running    1            21m
nginx-deployment-66b6c48dd5-5c4q6    1/1      Running    1            21m
nginx-deployment-66b6c48dd5-xs8nd    1/1      Running    1            21m
```

## What they do?

- create a set of resources, then control them as an unit
- **Operators** are an extension of the concept

# Custom Resources Definitions

- Define your own Kubernetes Resources
  - create new Kinds of resources
  - Handled as other resources

## Resource Handlers

- You need to write your own resource handler!
  - It responds to Kubernetes events
  - It interacts with Kubernetes APIs to perform operations



# Components of a CRD

- Group, Kind and short names:
  - Example: `nuvolaris.org`, `Sample`, `sam`
- Spec and Status
  - Versioned
  - defined as an OpenApi Schema:

```
type: object
properties:
  spec:
    type: object
```

# Defining a CRD (1/2)

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: samples.nuvolaris.org
spec:
  scope: Namespaced           # Namespaced or Cluster wide
  group: nuvolaris.org        # Group (resources are grouped)
  names:
    kind: Sample              # Kind (and its names)
    plural: samples
    singular: sample
    shortNames:
      - sam
```

## Defining a CRD (2/2)

```
versions:
  - name: v1
    served: true
    storage: true
    subresources: { status: { } }
    schema:
      openAPIV3Schema:
        type: object
        properties:
          spec:
            type: object
            x-kubernetes-preserve-unknown-fields: true
          status:
            type: object
            x-kubernetes-preserve-unknown-fields: true
```

# Instance

```
apiVersion: nuvolaris.org/v1
kind: Sample
metadata:
  name: obj
spec:
  count: 2
```

# Demo

```
cat demo-crd.yaml
kubectl apply -f demo-crd.yaml
kubectl get crd
cat demo-obj.yaml
kubectl apply -f demo-obj.yaml
kubectl get samples
kubectl delete sample obj
```

# Kustomize

# Interacting with Kubernetes

- Resources needs to be adapted to various cases
  - there are literally tens of solutions for this problem
  - `helm` is very used, but...
    - templating YAML is not a great idea!
    - lot of code just to manage whitespaces!
- We are going to use `kustomize`
  - part of `kubectl`
  - much easier development and debug

# About `kustomize`

- Originally a separate tool, now part of `kubectl`
  - It works "customizing" sets of descriptors with rules
  - support many ways of *patching* the JSON/YAML
  - **NO TEMPLATING** (huge win over `helm` !)
- You simply do `kubectl apply -k <folder>`
  - It will search for `kustomization.yaml`
  - It will produce the output sent to Kubernetes
- Debug the output without applying with:  
`kubectl kustomize <folder>`



# Simple `kustomization1.yaml` with patch

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
resources:
- demo-deployment.yaml
patches:
- path: patch.yaml
```

- put it in a folder `deploy` and `apply -k deploy`

# Sample patch of a Deployment

- We want to change the replica count

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: demo-deploy
spec:
  replicas: 2
```

- Intuitively, provide enough context to locate the descriptor
- Provide the replaced fields


# Kustomize example

```
cat demo-deployment.yaml | grep replica
# create a customization
mkdir deploy
cp demo-deployment.yaml deploy
cp patch.yaml deploy
cp kustomization.yaml deploy
# kustomization deploy
kubectl kustomize deploy | grep replica
kubectl apply -k deploy
kubectl get po
```

# Conclusion

# What is next?

- Setup Python and Kopf
- Authentication
- Reacting to Events
- Invoking Kubectl and Kustomize
- Implementing the Operator



## Nuvolaris Components and Technologies

# Contributing to Nuvolaris

## Before sending a Pull Request you need:

- Add **Apache License** headers to each file:
- The simplest way:  
`license-eye header fix`
- There is a check for each Pull Request

**Regular contributors need to sign the Apache ICLA**